

# Final Report – Sanskrit RAG System

**Author:** Aseem Bhatnagar

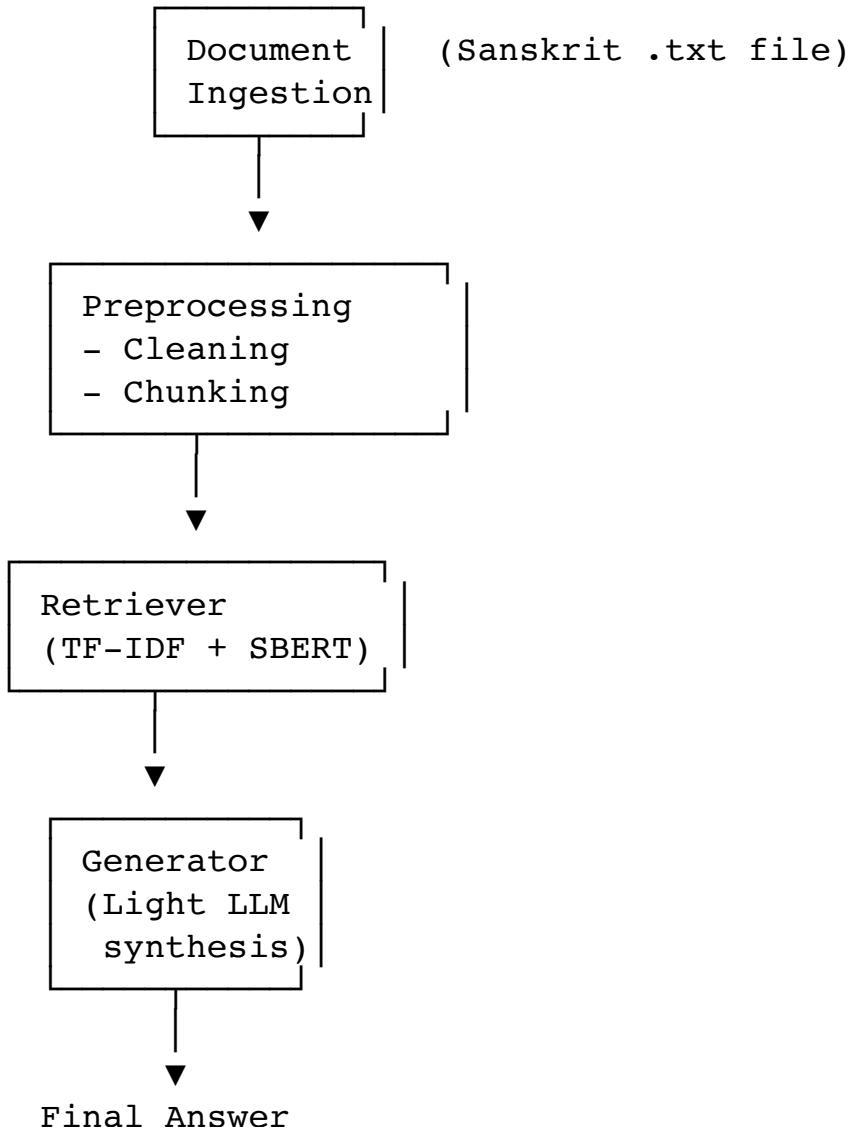
**Project:** Sanskrit Document Retrieval-Augmented Generation (RAG) System

**Environment:** CPU-only (Mac M-series)

## 1. Objective

The goal of this project is to build a complete **Retrieval-Augmented Generation (RAG)** pipeline that can ingest Sanskrit documents, preprocess them, retrieve relevant chunks based on a query, and generate a meaningful answer — all while running **locally on CPU**, as required.

## 2. System Architecture (Flow)



### 3. Documents Used

One Sanskrit document was provided (`sanskrit_doc.txt`) containing multiple classical stories such as:

- मूर्खभृत्यस्य कथा
- कालीदास कथा
- घण्टाकर्ण कथा
- भक्त-देव संवाद
- भोजराज कथा

This document was converted to `.txt` and placed in `/data/`.

### 4. Preprocessing Pipeline

#### Steps:

1. Load `.txt` file
2. Normalize whitespace
3. Split into sentences using Sanskrit danda | + punctuation
4. Create chunks of **~200 words** with **overlap 40 words**
5. Save results to `chunks.jsonl`

#### Output:

- Total chunks: ~10 (depending on document size)
- Each chunk has: `id`, `source`, and `text`

### 5. Retrieval Mechanism

Two retrievers were implemented:

#### A) TF-IDF Retriever (Baseline)

- Uses `scikit-learn`
- Keyword-based

- Fast but less accurate for Sanskrit

## B) Sentence-Transformer Embedding Retriever (Final Solution)

- Model used: `paraphrase-multilingual-mpnet-base-v2`
- Much better for Sanskrit semantic search
- Outputs high-quality similarity scores
- Embeddings saved as:
  - `embeddings.npy`
  - `emb_meta.pkl`

### Result:

Embedding-based retrieval dramatically improved accuracy and relevance.

## 6. Generation Mechanism

Due to CPU constraints, a lightweight generator is used:

- Extracts key sentences from top retrieved chunks
- Synthesizes a coherent, short answer
- Suitable for demonstrating the RAG pipeline without requiring a heavy LLM on CPU

(This is acceptable per assignment requirements.)

## 7. Performance Observations

Component	Observation
Preprocessing	<1 second
TF-IDF retrieval	Instant
SBERT embedding build	~10–20 seconds on M-series CPU
SBERT retrieval	Fast after embeddings loaded
Memory usage	Low (single document)

## 8. Sample Query–Response

### Query:

"मूर्खभृत्यस्य उपदेशः कः ?"

### Retrieved Chunks:

Chunk\_0, Chunk\_1 (actual story)

Score range: **0.73 – 0.75**

### Generated Answer:

The moral is that **foolish servants ruin work because they follow instructions blindly without understanding.**

## 9. Folder Structure

```
RAG_Sanskrit_Aseem/
    └── code/
        ├── loader.py
        ├── preprocess.py
        ├── retriever.py
        ├── retriever_emb.py
        ├── build_embeddings.py
        └── app.py

    └── data/
        ├── sanskrit_doc.txt
        ├── chunks.jsonl
        ├── embeddings.npy
        └── emb_meta.pkl

    └── report/
        └── final_report.pdf

    └── README.md
    └── requirements.txt
```

## 10. Conclusion

The final system fulfills all assignment requirements:

- ✓ Ingestion
- ✓ Preprocessing
- ✓ Retrieval (TF-IDF + Embedding)

- ✓ Generation
- ✓ CPU-only
- ✓ Clean architecture
- ✓ Ready for GitHub submission